

Entities

1. Users (userID, email (UNIQUE), userPassword, userBalance, userSignedUpAt)
 - Assumptions:
 - A user record represents a person using the app
 - userEmail must be unique so that no two users share the same login
 - userBalance is initialized with the amount the user deposits at sign-up
 - userSignedUpAt is automatically recorded at registration
 - A user's data persists even if they leave all groups
 - Why this is an entity:
 - Users are the main actors in the system and interact with multiple parts of the app (trading, groups, watchlists)
 - They exist independently of other entities and are referenced by many relationships
2. Groups (groupID, groupName, groupBalance):
 - Assumptions:
 - A group record represents a trading group mutually formed by users
 - A user can make multiple groups
 - groupName need not be unique
 - groupBalance is initialized with the collective balances of all users' individual balances
 - groupCreatedAt is automatically recorded when a group is made
 - groupCreatedBy references the user who created the group
 - Why this is an entity:
 - Groups exist independently of users and can include multiple users
 - Storing them separately avoids repeating group information for every user and supports group features like shared portfolios or discussions
3. Tickers (symbol, name, assetType):
 - Assumptions:
 - A ticker record represents a tradable symbol (e.g., "AAPL", "BTC", etc.)
 - symbol uniquely identifies each ticker
 - assetType classifies the ticker into common types (e.g., stocks, crypto, etc.)
 - Why this is an entity:

- Tickers are referenced across transactions, news, and price data
 - Keeping them as their own entity avoids redundancy and provides a single source of truth for all tradable instruments
4. Price-Bars (symbol, time, open, high, low, close, volume, source):
- Assumptions:
 - A price bar record represents market data for a specific ticker at a specific time
 - `symbol` and `time` combined uniquely identifies a record
 - `source` shows whether data is real or simulated
 - Market prices change frequently, leading to many records per ticker
 - Why this is an entity:
 - Price data forms a repeating time series for each ticker
 - Storing it separately allows the system to efficiently manage large amounts of market data without duplicating ticker information
5. News-Articles (articleID, `publishedAt`, `source`, `title`, `URL`, `sentiment`, `impactTags`):
- Assumptions:
 - A news article record represents an news article
 - `source`, `title`, and `URL` describe where the news article came from and what it is about
 - `sentiment` and `impactTags` are optional fields describing potential market effects
 - A single news article can mention multiple tickers
 - Why this is an entity:
 - News articles have their own metadata and aren't tied to a single ticker or user
 - Keeping them separate allows one news article to be linked to many tickers and easily reused or analyzed
6. Transactions (transactionID, `accountID`, `ticker`, `time`, `side`, `quantity`, `price`, `kind`, `status`, `requestedBy`, `approvedBy`):
- Assumptions:
 - A transaction record represents a single trade event
 - `transactionID` uniquely identifies each transaction
 - `accountID` links the trade to the account the trade originated from
 - Each trade involves one ticker and is requested by one user, as denoted with `requestedBy`
 - Some transactions may require approval from another user, as denoted with `approvedBy`
 - Transactions cannot be edited after its creation; in other words, they're treated as permanent records
 - Why this is an entity:

- Transactions are the core log of trading activity
- Since multiple accounts and users can each have many trades, storing them as their own entity allows for accurate tracking, auditing, and analysis

Relationships

1. Users ↔ Groups via Group-Memberships(role, joinedAt)
 - Cardinality: M:N (optional participation)
 - Users can belong to many groups
 - Groups can have many users
2. Users ↔ Groups via Created-By(createdAt, createdBy)
 - Cardinality: 1:M
 - Each group is created by exactly one user
 - A user can create many groups
3. Users ↔ Tickers via User-Watchlist
 - Cardinality: M:N (optional participation)
 - Each user can follow many tickers
 - Each ticker can be followed by many users
4. Users ↔ News-Articles via Users-News-Feed
 - Cardinality: M:N (optional participation)
 - Tracks which articles a user has seen/read
5. Users ↔ Transactions via User-Transactions-Requests(requestedBy)
 - Cardinality: 1:M
 - Each transaction request involves exactly one user
 - A user can request many transactions
6. Users ↔ Transactions via User-Transactions-Approvals(approvedBy)
 - Cardinality: 0..1:M
 - Each transaction can have either zero or one approver
 - Each user can approve many transactions
7. Groups ↔ Transactions via Groups-Transactions
 - Cardinality: 1:M
 - Each transaction involves exactly one group
 - A group can have many transactions
8. Tickers ↔ Transactions via Involves
 - Cardinality: 1:M
 - Each transaction involves exactly one ticker
 - A ticker can appear in many transactions
9. Tickers ↔ Price-Bars via HasPriceData
 - Cardinality: 1:M
 - Each ticker has many price bars

10. Tickers ↔ News-Articles via Mentions

- Cardinality: M:N (optional participation)
 - An article can mention many tickers
 - A ticker can appear in many articles

Normalizing

Fortunately enough, our schema is already in 3NF, seeing as how all non-key attributes depend on the whole key and there are no transitive dependencies

Relational Schema

1. Users(userID:INT [PK], userEmail:VARCHAR(255) UNIQUE, userPassword:VARCHAR(255), userBalance:DECIMAL(12,2), userSignedUpAt:DATETIME)
2. Grps(groupID:INT [PK], groupName:VARCHAR(255), groupBalance:DECIMAL(12,2), groupCreatedAt:DATETIME, groupCreatedBy:INT [FK to Users.userID])
3. Group-Memberships(userID: INT [FK to Users.userID],groupID: INT [FK to Groups.groupID], role: VARCHAR(50), joinedAt: DATETIME, PRIMARY KEY (userID, groupID))
4. Tickers(symbol: VARCHAR(10) [PK], name: VARCHAR(255), assetType: VARCHAR(50))
5. Price-Bars(ticker: VARCHAR(10) [FK to Tickers.symbol], time: DATETIME, open: DECIMAL(12,2), high: DECIMAL(12,2), low: DECIMAL(12,2), close: DECIMAL(12,2), volume: BIGINT, source: VARCHAR(20), PRIMARY KEY (ticker, time))
6. News-Articles(articleID: INT [PK], publishedAt: DATETIME, source: VARCHAR(100), title: VARCHAR(255), URL: VARCHAR(500), sentiment: VARCHAR(20), impactTags: VARCHAR(255))
7. Users-News-Feed(userID: INT [FK to Users.userID], articleID: INT [FK to News-Articles.articleID], seenAt: DATETIME, isRead: BOOLEAN, PRIMARY KEY (userID, articleID))
8. User-Watchlist(userID: INT [FK to Users.userID], ticker: VARCHAR(10) [FK to Tickers.symbol], addedAt: DATETIME, PRIMARY KEY (userID, ticker))
9. Transactions(transactionID: INT [PK], accountID: INT, groupID: INT [FK to Groups.groupID], ticker: VARCHAR(10) [FK to Tickers.symbol], time: DATETIME, side: VARCHAR(10), quantity: DECIMAL(12,4), price: DECIMAL(12,2), kind: VARCHAR(20), status: VARCHAR(20), requestedBy: INT [FK to Users.userID], approvedBy: INT [FK to Users.userID])